

Motivation
ooooo

Improve Interpretability and Transparency
ooooo

Safe Fail
ooooooooo

Verification of ML
ooooo

Assured Machine Learning

Xiaozhe Gu, Arvind Easwaran

School of Computer Science and Engineering
Energy Research Institute (ERI@N)

Nanyang Technological University, Singapore

June, 2018

Motivation
ooooo

Improve Interpretability and Transparency
ooooo

Safe Fail
ooooooooo

Verification of ML
ooooo

Outline

Motivation

Improve Interpretability and Transparency

Safe Fail

Verification of ML

Motivation
●○○○○

Improve Interpretability and Transparency
○○○○○

Safe Fail
○○○○○○○○○○

Verification of ML
○○○○○

Outline

Motivation

Improve Interpretability and Transparency

Safe Fail

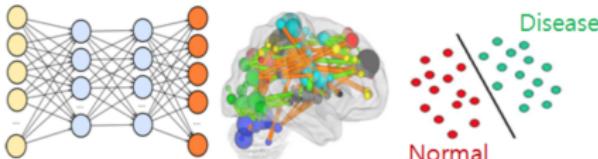
Verification of ML

Machine Learning Applications in Safety-Critical Environments

Machine learning algorithms are increasingly influencing our lives, and moving into safety-critical applications.

- ▶ Decision making in life-threatening conditions, e.g., ML based medical decision support systems.

Figure: ML Based Brain Disease Diagnosis[1]



Machine Learning Applications in Safety-Critical Environments

Machine learning algorithms are increasingly influencing our lives, and moving into **safety-critical** applications.

- ▶ Autonomous Robots, e.g., surgical robots, rescue robots, industrial robots, etc.

Figure: Surgical Robots[2]



Machine Learning Applications in Safety-Critical Environments

Machine learning algorithms are increasingly influencing our lives, and moving into safety-critical applications.

- ▶ Self-Driving Vehicles, e.g, autonomous shuttle.

Figure: Autonomous Shuttle



Machine Learning Applications in Safety-Critical Environments

Machine learning algorithms are increasingly influencing our lives, and moving into **safety-critical** applications.

- ▶ Autonomous Weapons

Figure: SGR-A1
(Source: wikipedia)



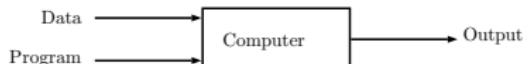
Figure: Source: keelvar.com



Traditional Programming versus Machine Learning



Figure: Traditional Programming



Traditional programming involves static program instructions **strictly specifying** what we need the computer to do.

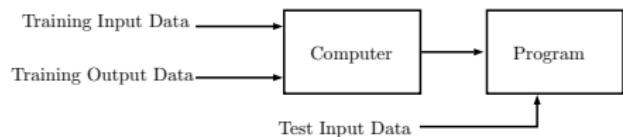
Traditional Programming versus Machine Learning

Figure: Traditional Programming



Traditional programming involves static program instructions strictly specifying what we need the computer to do.

Figure: Machine Learning



ML is a technique that enables computers to learn autonomously and to improve from experience without being explicitly programmed

Machine Learning Safety

Amodei et al. [5] refer to ML safety as “*mitigating risk in the context of unintended or harmful behaviour that may emerge from machine learning systems when we*”

- ▶ *specify the wrong objective function*,
 - ▶ *are not careful about the learning process*,
 - ▶ *or commit other machine learning-related implementation errors.*
-

Machine Learning Safety

Amodei et al. [5] refer to ML safety as “*mitigating risk in the context of unintended or harmful behaviour that may emerge from machine learning systems when we*”

- ▶ *specify the wrong objective function*,
 - ▶ *are not careful about the learning process*,
 - ▶ *or commit other machine learning-related implementation errors.*
-

Bostrom et al.[6] refer to safety as “*techniques that ensure that machine learning systems behave as intended*”.

Challenges to Safety Assurance

- ▶ **Non-transparency:** The reasoning behind some powerful ML models is not known.

Challenges to Safety Assurance

- ▶ **Non-transparency:** The reasoning behind some powerful ML models is not known.
- ▶ **Unmodeled phenomena:** It is impossible and not desired to model everything.

Challenges to Safety Assurance

- ▶ **Unmodeled phenomena:** It is impossible and not desired to model everything.
We train a model to recognise dog breeds , but are given a cat to classify.



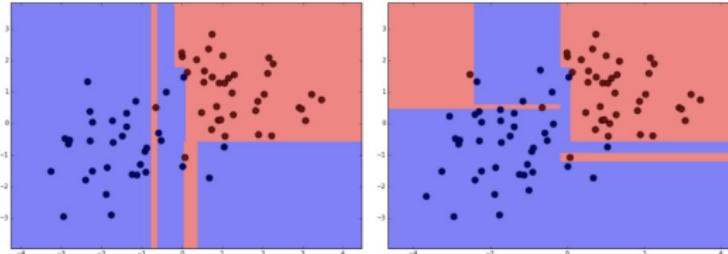
Challenges to Safety Assurance

- ▶ **Non-transparency**: The reasoning behind some powerful ML models is not known.
- ▶ **Unmodeled phenomena**: It is impossible and not desired to model everything.
- ▶ **Instability**: A small change in the training process may produce a different result.

Challenges to Safety Assurance

- ▶ **Non-transparency:** The reasoning behind some powerful ML models is not known.
- ▶ **Unmodeled phenomena:** It is impossible and not desired to model everything.
- ▶ **Instability:** A small change in the training process may produce a different result.

Figure: Little variation in training procedure produce different classification rule



Challenges to Safety Assurance

- ▶ **Non-transparency**: The reasoning behind some powerful ML models is not known.
- ▶ **Unmodeled phenomena**: It is impossible and not desired to model everything.
- ▶ **Instability**: A small change in the training process may produce a different result.
- ▶ **Difficulty in verification**: Formal verification of ML components is a difficult, and somewhat ill-posed problem.

Motivation
ooooo

Improve Interpretability and Transparency
●oooo

Safe Fail
ooooooooo

Verification of ML
ooooo

Outline

Motivation

Improve Interpretability and Transparency

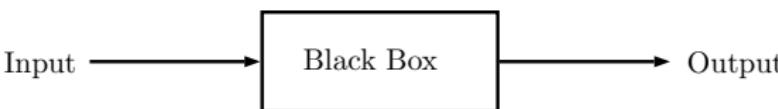
Safe Fail

Verification of ML

Non-transparency

Many machine learning models (e.g., deep neural network) behave mostly as black boxes.

Figure: Black Box



Understanding the reasons behind model/prediction is, however, quite important in assessing trust without which the users may not deploy/use it.

Interpretable Models

Very common model types for interpretable models are:

- ▶ Linear regression model

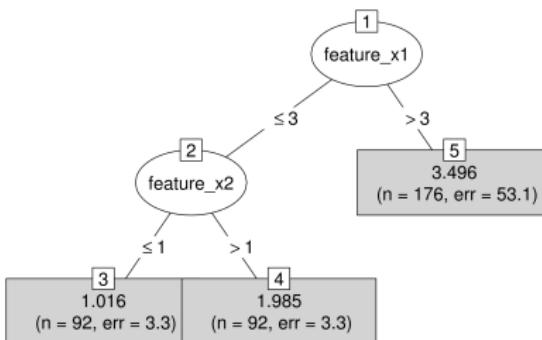
$$\begin{aligned}y &= \mathbf{w}^T \cdot \mathbf{x} + \epsilon \\&= w_0 + \sum_{i=1}^m w_i x_i + \epsilon\end{aligned}$$

Interpretable Models

Very common model types for interpretable models are:

- ▶ Decision trees.

Figure: Regression Tree



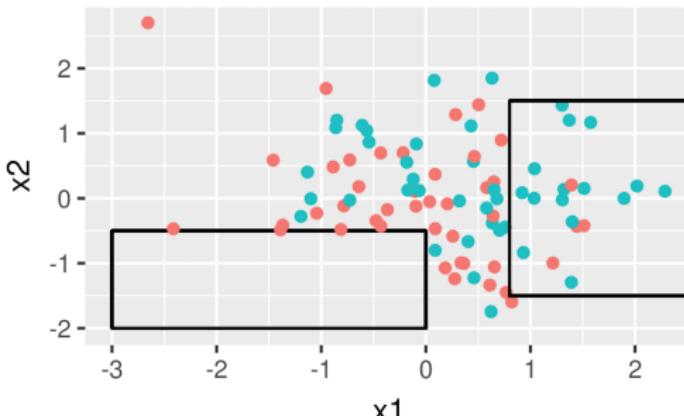
Interpretable Models

Very common model types for interpretable models are:

- ▶ Decision Rules:

If $-3 \leq x_1 \leq 0$ and $-2 \leq x_2 \leq -0.5$: then $y = +1$

Figure: Decision Rule



Interpretable Models

Very common model types for interpretable models are:

- ▶ Naive Bayes classifier

$$\begin{aligned} P(y|x_1 \dots x_d) &= \frac{P(y)P(x_1 \dots x_d|y)}{P(x_1 \dots x_d)} \\ &= \frac{P(y) \prod_{i=1}^d P(x_i|y)}{P(x_1 \dots x_d)} \\ &\propto P(y) \prod_{i=1}^d P(x_i|y) \\ \Rightarrow y &= \arg \max_y P(y) \prod_{i=1}^d P(x_i|y) \end{aligned}$$

Interpretable Models

Very common model types for interpretable models are:

- ▶ k-Nearest Neighbors

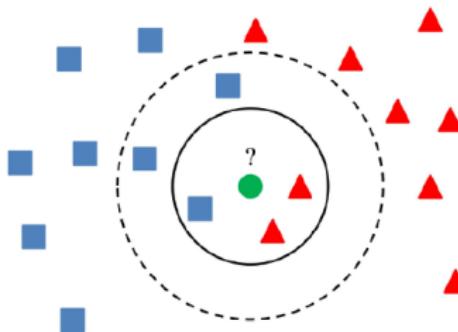
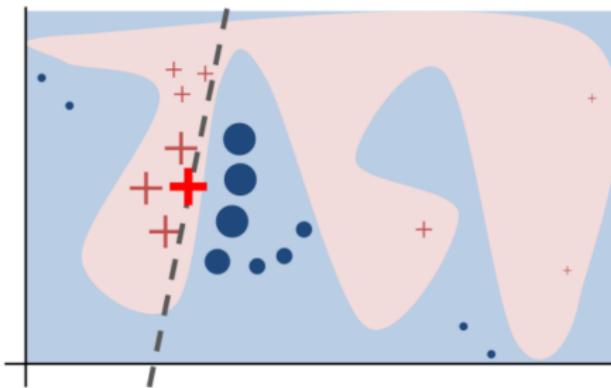


Figure: KNN (Source:wikipedia)

Explaining the Predictions

Explain the prediction of a classifier by approximating it locally with an interpretable model [3].

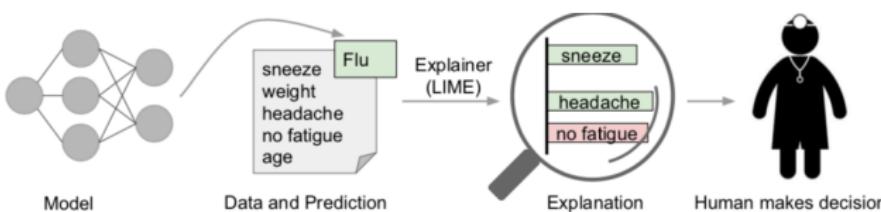
Figure: The black-box model's complex decision represented by the blue/pink background is approximated locally by a linear model [3].



Explaining the Predictions

With the explanation of a prediction, a doctor can make an informed decision about whether to trust the model's prediction.

Figure: Explaining individual predictions [3]



Global Surrogate Models

A global surrogate model is an interpretable model that is trained to approximate a black box model

How well the surrogate replicates the black box model?

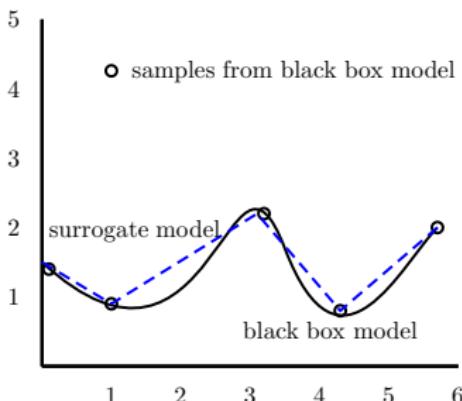
$$\text{Metric: } 1 - \frac{\sum_{i=1}^n (\hat{y}_i^* - \hat{y}_i)^2}{\sum_{i=1}^n (\hat{y}_i^* - \bar{\hat{y}}_i)}$$

where \hat{y}_i^* and \hat{y}_i is the prediction of the surrogate model and the black box model respectively.

Global Surrogate Models

A global surrogate model is an interpretable model that is trained to approximate a black box model

Figure: Linear Surrogate Model



Motivation
ooooo

Improve Interpretability and Transparency
ooooo

Safe Fail
●ooooooooo

Verification of ML
ooooo

Outline

Motivation

Improve Interpretability and Transparency

Safe Fail

Verification of ML

Learning to Reject

A technique used in machine learning when predictions cannot be given confidently is the **reject option** [4].

$$f(x_i) = \text{rejection if } g(f, x_i) \leq \sigma$$

where $g(f, x_i)$ measures the confidence level of function f's prediction for x_i , and σ is a threshold.

Learning to Reject

A technique used in machine learning when predictions cannot be given confidently is the **reject option** [4].

$$f(x_i) = \text{rejection if } g(f, x_i) \leq \sigma$$

where $g(f, x_i)$ measures the confidence level of function f 's prediction for x_i , and σ is a threshold.

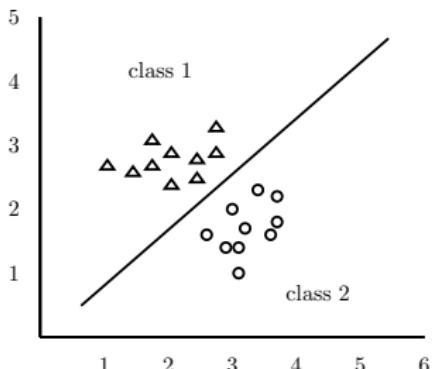
When the model selects the reject option, a **human operator** can intervene and provide a manual prediction.

Confidence: Distance from Decision Boundary

In classification problems, classifier implicitly assumes that distance from the decision boundary is inversely related to confidence [7].



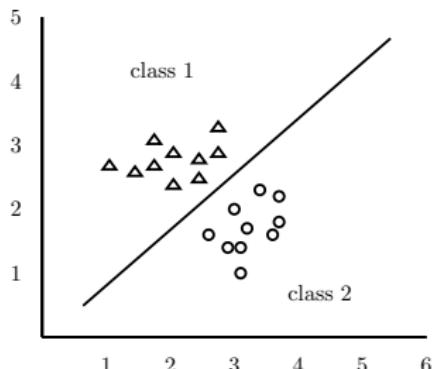
Figure: Decision Boundary



Confidence: Distance from Decision Boundary

In classification problems, classifier implicitly assumes that distance from the decision boundary is inversely related to confidence [7].

Figure: Decision Boundary



This is reasonable *to some degree* because the decision boundary is located where there is a large overlap in likelihood functions.

Confidence: Votes

In ensemble classification, the overall decision \hat{y} is based on the average classification of the base classifiers $\hat{y}_i(\cdot)$,

$$\phi(\mathbf{x}) = \frac{1}{m} \sum_{i=1}^m \hat{y}_i(\mathbf{x})$$

Confidence: Votes

In ensemble classification, the overall decision \hat{y} is based on the average classification of the base classifiers $\hat{y}_i(\cdot)$,

$$\phi(\mathbf{x}) = \frac{1}{m} \sum_{i=1}^m \hat{y}_i(\mathbf{x})$$

The reject option of ensemble binary classification is based on the votes of the base classifiers[7]

$$\hat{y}(x) = \begin{cases} -1 & \text{if } \phi(\mathbf{x}) \leq t_1 \\ \text{reject, if } \phi(x) \in (t_1, t_2) \\ 1 & \text{if } \phi(x) \geq t_2 \end{cases}$$

Confidence: Softmax Output

In neural network, predictive probabilities obtained at the end of the softmax output are often erroneously interpreted as model confidence [12].

$$P(y = j|\mathbf{x}) = \frac{e^{\mathbf{x}^\top \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^\top \mathbf{w}_k}}$$

Confidence: Softmax Output

In neural network, predictive probabilities obtained at the end of the softmax output are often erroneously interpreted as model confidence [12].

$$P(y = j|\mathbf{x}) = \frac{e^{\mathbf{x}^\top \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^\top \mathbf{w}_k}}$$

However, a model can be uncertain in its predictions even with a high softmax output [12].

Prediction Uncertainty: Gaussian Process

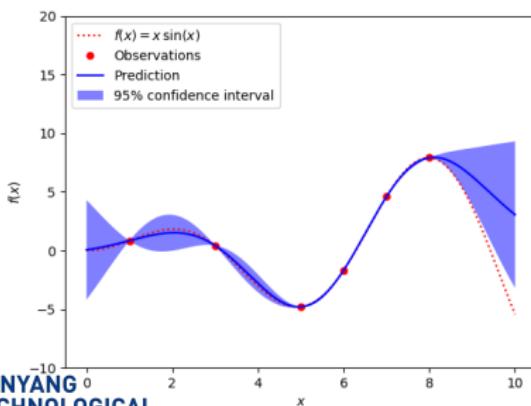
A Gaussian process defines a distribution over functions $p(f)$ where f is a function mapping some input space $\mathcal{X} \rightarrow \mathbb{R}$.

$$p(\mathbf{f}) \sim \mathcal{N}(0, \Sigma) = \left(\frac{1}{2\pi|\Sigma|} \right)^{D/2} \exp(-1/2\mathbf{f}^T \Sigma^{-1} \mathbf{f})$$

Prediction Uncertainty: Gaussian Process

A Gaussian process defines a distribution over functions $p(f)$ where f is a function mapping some input space $\mathcal{X} \rightarrow \mathbb{R}$.

Figure: One-dimensional Regression



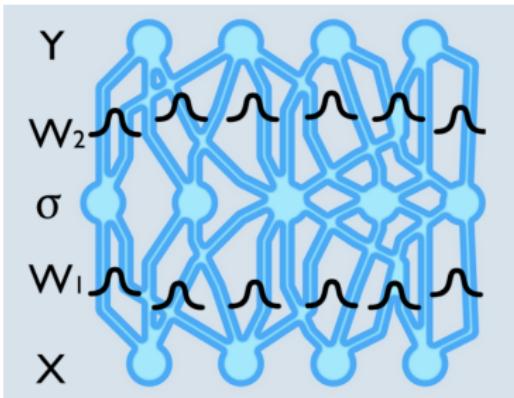
Prediction for input x with larger standard deviation implicates lower confidence. Thus, we can reject to predict for input x with larger standard deviation.

Prediction Uncertainty: Bayesian Neural Network

A Bayesian neural network is a neural network with a prior distribution on its weights [13], e.g.,

$$p(\mathbf{w}) \sim \mathcal{N}(0, \sigma^2 I)$$

Figure: Prior distribution on weight [13]



Prediction Uncertainty: Bayesian Neural Network

A Bayesian neural network is a neural network with a prior distribution on its weights [13], e.g.,

$$p(\mathbf{w}) \sim \mathcal{N}(0, \sigma^2 I)$$

posterior:

$$P(\mathbf{w}|Y, X) \propto P(Y|\mathbf{w}, X)p(\mathbf{w})$$

Prediction Uncertainty: Bayesian Neural Network

A Bayesian neural network is a neural network with a prior distribution on its weights [13], e.g.,

$$p(\mathbf{w}) \sim \mathcal{N}(0, \sigma^2 I)$$

posterior:

$$P(\mathbf{w}|Y, X) \propto P(Y|\mathbf{w}, X)p(\mathbf{w})$$

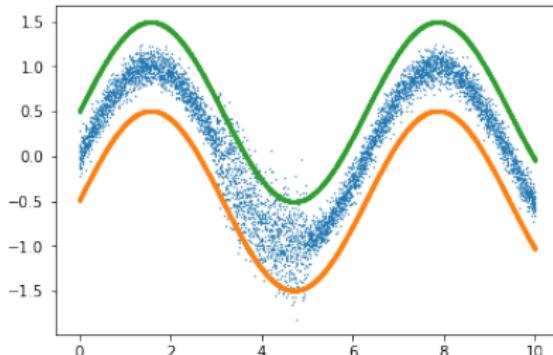
Inference:

$$p(\mathbf{y}_*|\mathbf{x}_*, X, Y) = \int p(\mathbf{y}_*|\mathbf{x}_*, \mathbf{w})p(\mathbf{w}|X, Y)d\mathbf{w}$$

Prediction with Confidence Intervals: Conformal Prediction

Conformal prediction [8] gives us valid bounds $[f(\mathbf{x}) - \epsilon_1, f(\mathbf{x}) + \epsilon_2]$ for any model such that the prediction region contains the true output with probability α .

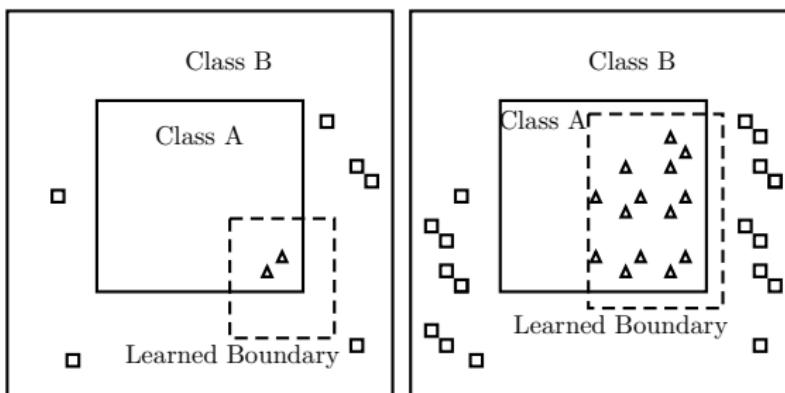
Figure: The Bound with 95% Confidence



Reject Unmodeled Scenario

It is impossible for ML to model everything. Thus, if the model recognizes that an instance is from the input space that is not well trained, it could reject to predict.

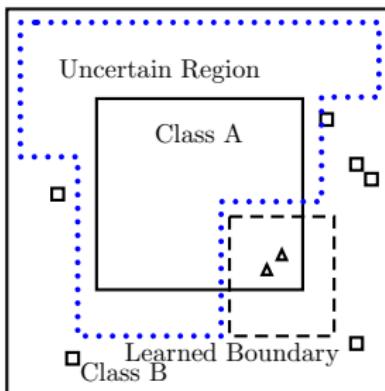
Figure: Wrong Decision Boundary due to Lack of Data



Reject Unmodeled Scenario

It is impossible for ML to model everything. Thus, if the model recognizes that an instance is from the input space that is not well trained, it could reject to predict.

Figure: Region with Low-confidence



Motivation
ooooo

Improve Interpretability and Transparency
ooooo

Safe Fail
ooooooooo

Verification of ML
●oooo

Outline

Motivation

Improve Interpretability and Transparency

Safe Fail

Verification of ML

Challenges for Verified ML

The formal verification of ML components is a difficult, and somewhat ill-posed problem due to the complexity of the ML algorithms, large feature spaces, etc. [9].

Challenges for Verified ML

Challenges to achieving formally-verified ML-based systems:

- ▶ Environment Modeling:

It may be impossible even to precisely define all the variables (features) of the environment that must be modeled, let alone to model all possible behaviors of the environment.



Challenges for Verified ML

Challenges to achieving formally-verified ML-based systems:

► Specification:

How to specify desired and undesired properties of ML components. For example, how to formally specify “a dog” ?



Challenges for Verified ML

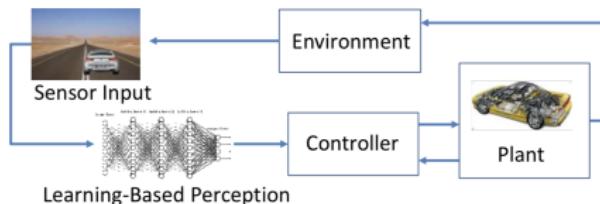
Challenges to achieving formally-verified ML-based systems:

- ▶ System Modeling:

Unlike traditional applications of formal verification, ML based system evolves as it encounters new data and new situations. Thus, it is difficult to reuse the previous safety verification results.

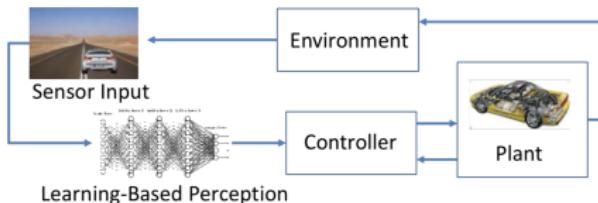
Go System-Level

- ▶ Instead of verifying the ML component directly,
- ▶ Formally specify the end-to-end behavior of the system and verify the system containing the ML component [10].



Go System-Level

- ▶ Instead of verifying the ML component directly,
- ▶ Formally specify the end-to-end behavior of the system and verify the system containing the ML component [10].

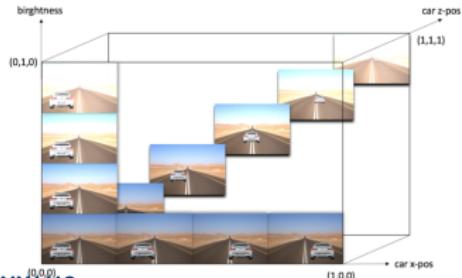


Specification: Always $\text{dist}(\text{ego vehicle}, \text{env object}) > \Delta$

Approximate Model in Abstract Space

Abstract Feature Space: Instead of the high-dimensional input space, explore realistic and meaningful simplifications [10].

Figure: The abstract space A with the three dimensions



Approximate Model in Abstract Space

Abstract Feature Space: Instead of the high-dimensional input space, explore realistic and meaningful simplifications [10].

A simpler approximate function \hat{f} of the original model f in the abstract domain is analyzed.

Figure: The abstract space A with the three dimensions

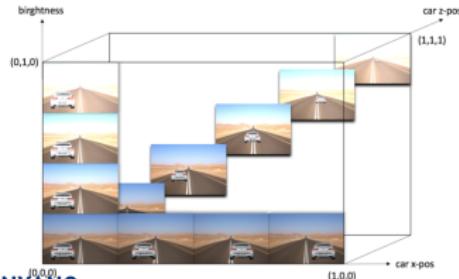
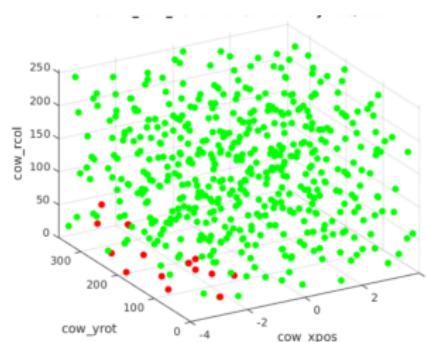


Figure: Misclassifying Elements



Reduce the Problem

- ▶ Reduce the input set \mathcal{X} of n dimension into a finite number of regularized subregions. For example, partition an input set into hyper-rectangles [11].

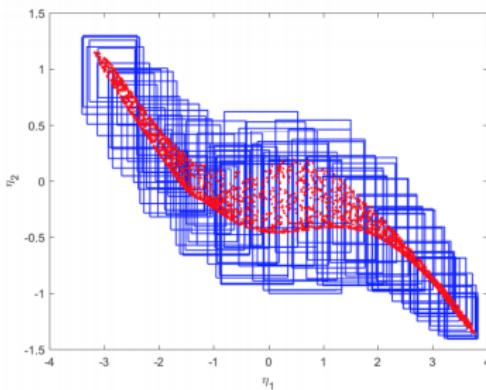
$$\mathcal{P}_i \leftarrow \mathcal{I}_{1,m_1} \times \mathcal{I}_{2,m_2} \times \dots \times \mathcal{I}_{n,m_n}$$

where $I_{i,m_i} = [\eta_{m_i-1}, \eta_{m_i}]$

Reduce the Problem

- ▶ For each hyper-rectangle input set, over-approximate the output set by a hyper-rectangle $[\underline{\phi}_j, \bar{\phi}_j]$

Figure: Blue rectangles: the over-approximated output set, red spots: 5000 random outputs located in the estimated output set [11]



- ❑ <http://mlcenter.postech.ac.kr/healthcare>
- ❑ <https://www.wired.com/2015/03/google-robot-surgery/>
- ❑ Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin. Why should i trust you?: Explaining the predictions of any classifier. Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. ACM, 2016.
- ❑ Bartlett P L, Wegkamp M H. Classification with a reject option using a hinge loss[J]. Journal of Machine Learning Research, 2008, 9(Aug): 1823-1840.
- ❑ Amodei, Dario, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mane. 2016. Concrete Problems in AI Safety.

-  Bostrom, N., Dafoe, A., and Flynn, C. 2016. Policy Desiderata in the Development of Machine Superintelligence
-  K. R. Varshney, R. J. Prenger, T. L. Marlatt, B. Y. Chen, and W. G. Hanley, Practical ensemble classification error bounds for different operating points, IEEE Transactions on Knowledge and Data Engineering, vol. 25, no. 11, pp. 2590-2601, Nov. 2013.
-  Vovk V, Gammerman A, Shafer G. Conformal prediction[M]. Springer US, 2005
-  S. A. Seshia, D. Sadigh, and S. S. Sastry. Towards verified artificial intelligence. CoRR, abs/1606.08514, 2016.
-  Tommaso Dreossi, Alexandre Donze, Sanjit A. Seshia, Compositional Falsification of Cyber-Physical Systems with Machine Learning Components, Preprint/

-  Weiming Xiang and Taylor T. Johnson, Reachability Analysis and Safety Verification for Neural Network Control Systems, Preprint.
-  Gal Y, Ghahramani Z. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning[C]//international conference on machine learning. 2016: 1050-1059.
-  Neal, R. M. (2012). Bayesian learning for neural networks (Vol. 118). Springer Science & Business Media.